

MZT

Laboratorium 1

Przemysław Kleszcz

Nr albumu: 124624

Platforma testowa

Procesor	Intel Core i5-7440HQ, 2.8GHz, 6MB Cache
RAM	DDR4 - 16 GB
System operacyjny	Microsoft Windows 10 Pro
Środowisko programistyczne	Visual Studio Professional 2017 v15.5.7
Środowisko uruchomieniowe	.NET Framework 4.6.01055

Zadanie 1

```
void matvec_opt_1(double* a, double* x, double* y, int n)
/*=====
Usuniecie skokow w danych
=====*/
{
    int i, j = 0, k = 0;
    memset(y, 0, n * sizeof(double));
    double register r;
    for (j = 0; j < n; j++) {
        r = x[j];
        for (i = 0; i < n; i++) {
            y[i] += a[k++] * r;
        }
    }
}
```

```

void matvec_opt_2(double* a, double* x, double* y, int n)
/*=====
Rozwijanie petli
=====*/
{
    int rest = n % 8;
    int i = 0, j = 0, ij = 0;
    double register r = 0;

    memset(y, 0, n * sizeof(double));

    for (j = 0; j < n; ++j) {
        r = x[j];
        for (i = 0; i < n - rest; i += 8) {
            y[i] += a[ij] * r;
            y[i + 1] += a[ij + 1] * r;
            y[i + 2] += a[ij + 2] * r;
            y[i + 3] += a[ij + 3] * r;
            y[i + 4] += a[ij + 4] * r;
            y[i + 5] += a[ij + 5] * r;
            y[i + 6] += a[ij + 6] * r;
            y[i + 7] += a[ij + 7] * r;
            ij += 8;
        }

        for (; i < n; ++i) {
            y[i] += a[ij++] * r;
        }
    }
}

```

D:\Projects\MZT\Z1\x64\Release\MZT_L1.exe

```

Program Mat_Vect: performance y = y +A*x
x64 plarform
RELEASE version
Input dimension
10000
start
non-optimized code:
CPU time: 0.919 s
optim_1:
CPU time: 0.17 s
OK
optim_2:
CPU time: 0.17 s
OK
optim_3:
CPU time: 0.064 s
OK
optim_4:
CPU time: 0.074 s
OK
optim_5: AVX
CPU time: 0.066 s
OK
optim_6: FMA
CPU time: 0.072 s
OK
Press any key to continue . . .

```

Zadanie 2

```
void matvec_XMM(double* a, double* x, double* y, int n, int lb)
{
    int i, j;
    __m128d rx0, ra0, ry0;
    double *ptr_x, *ptr_a;
    __declspec(align(16)) double tmp[2];
    memset((void *)y, 0, n * sizeof(double));
    ptr_a = a;
    for (i = 0; i < n; i++)
    {
        ry0 = _mm_setzero_pd();
        ptr_x = x;
        for (j = 0; j < n; j += 2, ptr_a += 2, ptr_x += 2)
        {
            rx0 = _mm_load_pd(ptr_x);
            ra0 = _mm_load_pd(ptr_a);
            ra0 = _mm_mul_pd(ra0, rx0);
            ry0 = _mm_add_pd(ry0, ra0);
        }
        _mm_store_pd(tmp, ry0);
        y[i] = tmp[0] + tmp[1];
    }
}

void matvec_YMM(double* a, double* x, double* y, int n, int lb)
{
    int i, j;
    __m256d rx0, ra0, ra1, ra2, ra3, ry0, ry1, ry2, ry3;
    double *ptr_x, *ptr_a;
    __declspec(align(16)) double tmp0[4], tmp1[4], tmp2[4], tmp3[4];
    memset((void *)y, 0, n * sizeof(double));
    ptr_a = a;
    for (i = 0; i < n; i += 4)
    {
        ry0 = ry1 = ry2 = ry3 = _mm256_setzero_pd();
        ptr_x = x;
        for (j = 0; j < n; j += 16)
        {
            _mm_prefetch((const char *) (ptr_x + 16), _MM_HINT_T0);
            _mm_prefetch((const char *) (ptr_x + 24), _MM_HINT_T0);
            _mm_prefetch((const char *) (ptr_a + 16), _MM_HINT_NTA);
            _mm_prefetch((const char *) (ptr_a + 24), _MM_HINT_NTA);
            //poprawic
            //-----0
            rx0 = _mm256_load_pd(ptr_x);
            ra0 = _mm256_load_pd(ptr_a);
            ra1 = _mm256_load_pd(ptr_a + n);
            ra2 = _mm256_load_pd(ptr_a + 2 * n);
            ra3 = _mm256_load_pd(ptr_a + 3 * n);
            ra0 = _mm256_mul_pd(ra0, rx0);
            ra1 = _mm256_mul_pd(ra1, rx0);
            ra2 = _mm256_mul_pd(ra2, rx0);
            ra3 = _mm256_mul_pd(ra3, rx0);
            ry0 = _mm256_add_pd(ry0, ra0);
            ry1 = _mm256_add_pd(ry1, ra1);
        }
    }
}
```

```

ry2 = _mm256_add_pd(ry2, ra2);
ry3 = _mm256_add_pd(ry3, ra3);
//-----256-----1
rx0 = _mm256_load_pd(ptr_x + 4);
ra0 = _mm256_load_pd(ptr_a + 4);
ra1 = _mm256_load_pd(ptr_a + n + 4);
ra2 = _mm256_load_pd(ptr_a + 2 * n + 4);
ra3 = _mm256_load_pd(ptr_a + 3 * n + 4);
ra0 = _mm256_mul_pd(ra0, rx0);
ra1 = _mm256_mul_pd(ra1, rx0);
ra2 = _mm256_mul_pd(ra2, rx0);
ra3 = _mm256_mul_pd(ra3, rx0);
ry0 = _mm256_add_pd(ry0, ra0);
ry1 = _mm256_add_pd(ry1, ra1);
ry2 = _mm256_add_pd(ry2, ra2);
ry3 = _mm256_add_pd(ry3, ra3);
//-----256-----2
rx0 = _mm256_load_pd(ptr_x + 8);
ra0 = _mm256_load_pd(ptr_a + 8);
ra1 = _mm256_load_pd(ptr_a + n + 8);
ra2 = _mm256_load_pd(ptr_a + 2 * n + 8);
ra3 = _mm256_load_pd(ptr_a + 3 * n + 8);
ra0 = _mm256_mul_pd(ra0, rx0);
ra1 = _mm256_mul_pd(ra1, rx0);
ra2 = _mm256_mul_pd(ra2, rx0);
ra3 = _mm256_mul_pd(ra3, rx0);
ry0 = _mm256_add_pd(ry0, ra0);
ry1 = _mm256_add_pd(ry1, ra1);
ry2 = _mm256_add_pd(ry2, ra2);
ry3 = _mm256_add_pd(ry3, ra3);
//-----256-----3
rx0 = _mm256_load_pd(ptr_x + 12);
ra0 = _mm256_load_pd(ptr_a + 12);
ra1 = _mm256_load_pd(ptr_a + n + 12);
ra2 = _mm256_load_pd(ptr_a + 2 * n + 12);
ra3 = _mm256_load_pd(ptr_a + 3 * n + 12);
ra0 = _mm256_mul_pd(ra0, rx0);
ra1 = _mm256_mul_pd(ra1, rx0);
ra2 = _mm256_mul_pd(ra2, rx0);
ra3 = _mm256_mul_pd(ra3, rx0);
ry0 = _mm256_add_pd(ry0, ra0);
ry1 = _mm256_add_pd(ry1, ra1);
ry2 = _mm256_add_pd(ry2, ra2);
ry3 = _mm256_add_pd(ry3, ra3);
ptr_a += 16;
ptr_x += 16;
}
ptr_a += 3 * n;
_mm256_store_pd(tmp0, ry0);
_mm256_store_pd(tmp1, ry1);
_mm256_store_pd(tmp2, ry2);
_mm256_store_pd(tmp3, ry3);
y[i] = tmp0[0] + tmp0[1] + tmp0[2] + tmp0[3];
y[i + 1] = tmp1[0] + tmp1[1] + tmp1[2] + tmp1[3];
y[i + 2] = tmp2[0] + tmp2[1] + tmp2[2] + tmp2[3];
y[i + 3] = tmp3[0] + tmp3[1] + tmp3[2] + tmp3[3];
}
}

```

D:\Projects\MZT\Z2\x64\Debug\MatVect_lab_vect_regist.exe

Program Mat_Vect: performance $y = y + A*x$

DEBUG version

Input dimension

10000

start

naive algorithm:

naive algorithm: 0.375 sec

OK

algorithm which uses XMM: 0.766 sec

OK

algorithm which uses YMM: 0.078 sec

OK

algorithm which uses FMA: 0 sec

Press any key to continue . . .